

# OS 2018

Homework 4: Memory Allocator Implementation

(Due date 01/03 23:59:59)

# Objective

- Understand how **malloc()** and **free()** work
- Understand how to **manage Heap**

# Requirements

# Requirements

1. Implement a **Memory Allocator Library** for user application (*slides 4-16*)
2. Write a **User Application** to test the memory allocator library (*slides 17,18*)
3. Please follow the **Input/output format** (*slides 19,20*)

# 1. Memory Allocator Library Requirements

- The library must provide the following **3 functions**, and you should follow the format in *next slide* :
  1. **hw\_malloc()** (*slides 6-13*)
  2. **hw\_free()** (*slides 14-16*)
  3. **hw\_get\_start\_brk()**
- Use **chunk** (*slides 7,8*), **bin** (*slides 9,10*) and **sbrk()** to **manage heap**
- Use **chunk** (*slides 7,8*) to **manage every mmap-allocated memory**

# Functions format

## 1. `void *hw_malloc(size_t bytes)`

- bytes : the required memory size in bytes
- $\text{Return} = \begin{cases} \text{the valid virtual address} \\ \text{(starting address of the data part)} & , \text{ if success} \\ \text{NULL} & , \text{ otherwise} \end{cases}$

## 2. `int hw_free(void *mem)`

- mem : starting address of the data part
- $\text{Return} = \begin{cases} 1 & , \text{ if success} \\ 0 & , \text{ otherwise} \end{cases}$

## 3. `void *hw_get_start_brk()`

- Return the starting address of the heap

# hw\_malloc( ) requirements

- Use **mmap\_threshold** to decide the memory allocate method. If **the allocated size(data size + chunk header size) > mmap\_threshold** , use mmap allocation method ; else use Heap allocation method.
  - **mmap\_threshold** is **initial : 32 KiB** (32 \* 1024)

## mmap allocation method :

- Use **mmap()** system call to allocate the space
- Use **chunk** (slides 7,8) to manage the allocated space

The allocate size = request size (data size) + chunk header size

- Use **mmap\_alloc\_list** (slide 11) to manage allocated mmap chunks

## Heap allocation method:

- Use **chunk** (slides 7,8) and **bin** (slides 9,10) to manage heap
- Should follow the rules of **Heap initialization** (slide 12)
- Should follow the rules of **Split** (slide 13)
- **The allocated size (data size + chunk header size) should be the best fit size**

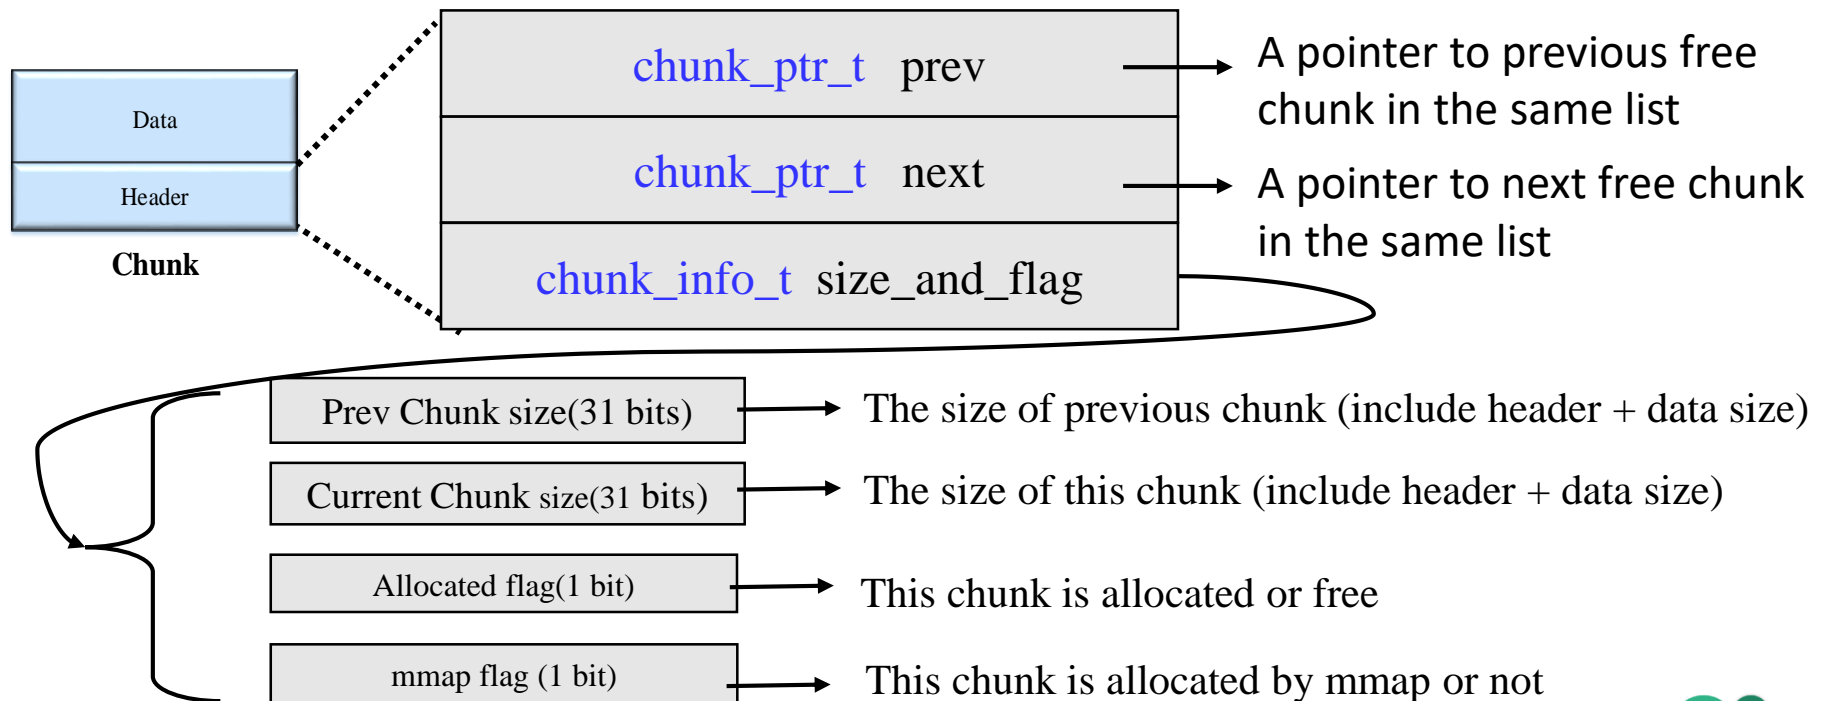
# Chunk requirements

- The continuous heap space is split into chunk(s) for management
- Each chunk contains two parts, header and data (*in next slide*)
  - Header (lower address)
  - Data (higher address), the actual memory space return to caller



# Chunk header format

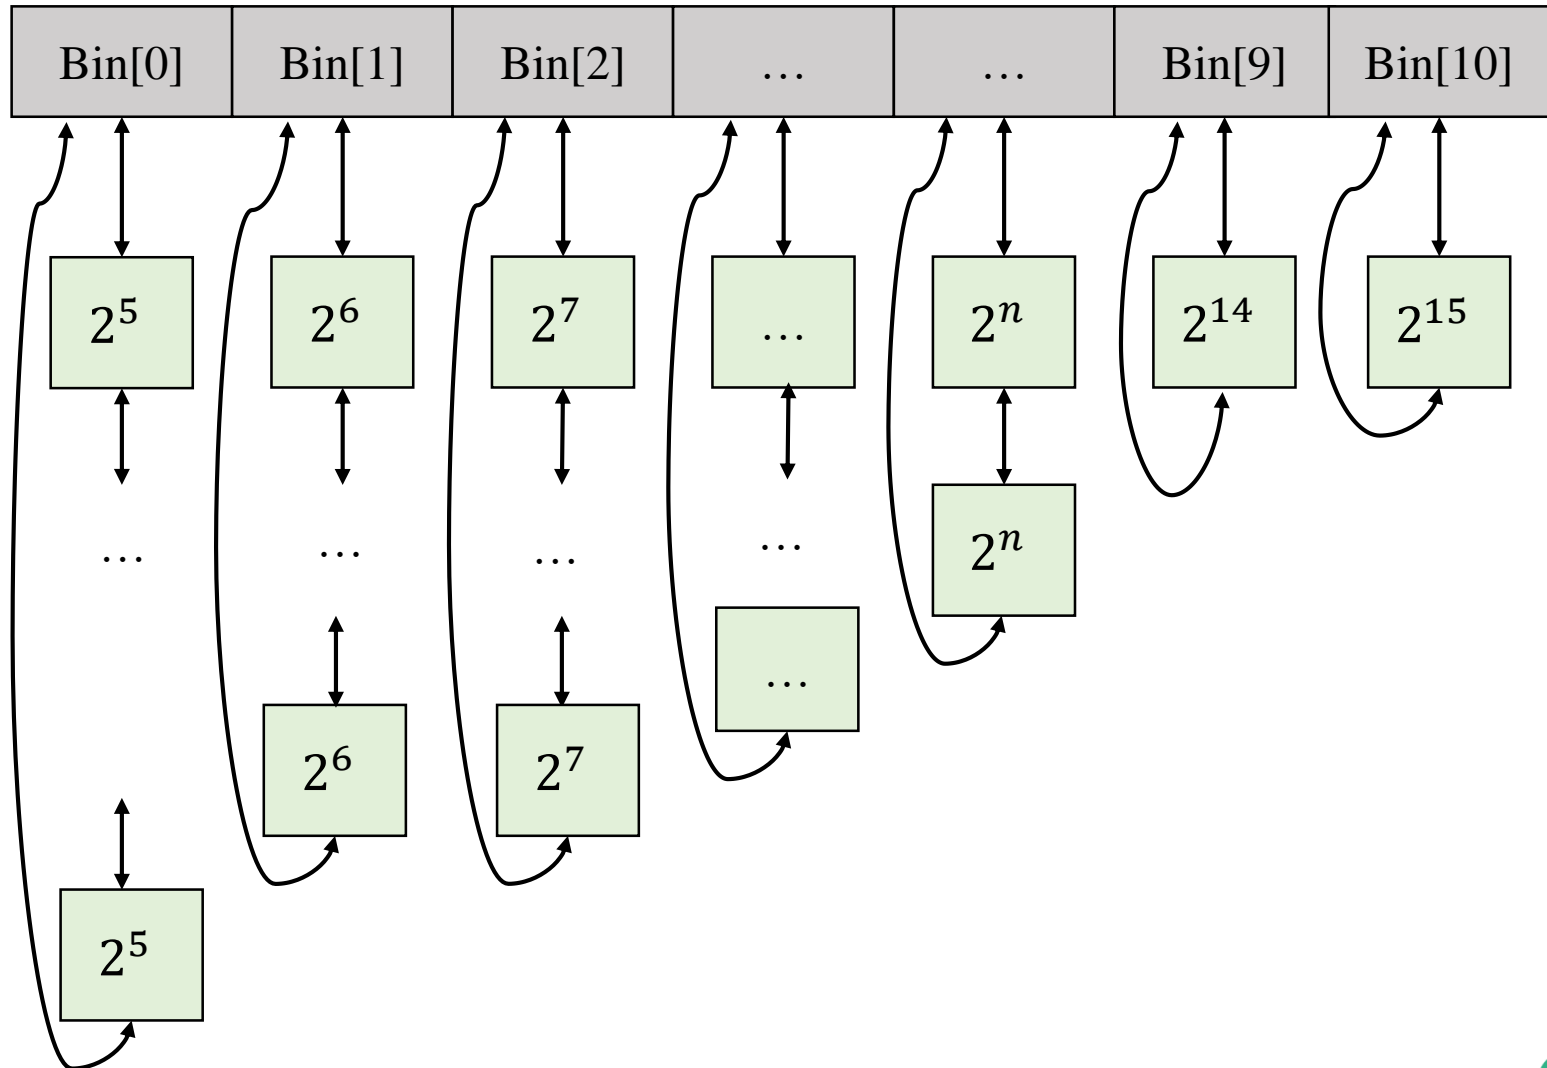
- Chunk header (24 bytes)
  - There are 3 members in the header
    - `chunk_ptr_t`, `chunk_size_t`, and `chunk_sizeandflag_t` can be defined by yourself, but each of them should be 8 bytes
    - `chunk_info_t` should include 4 information (Prev Chunk size, Current Chunk size, Allocates flag and mmap flag)



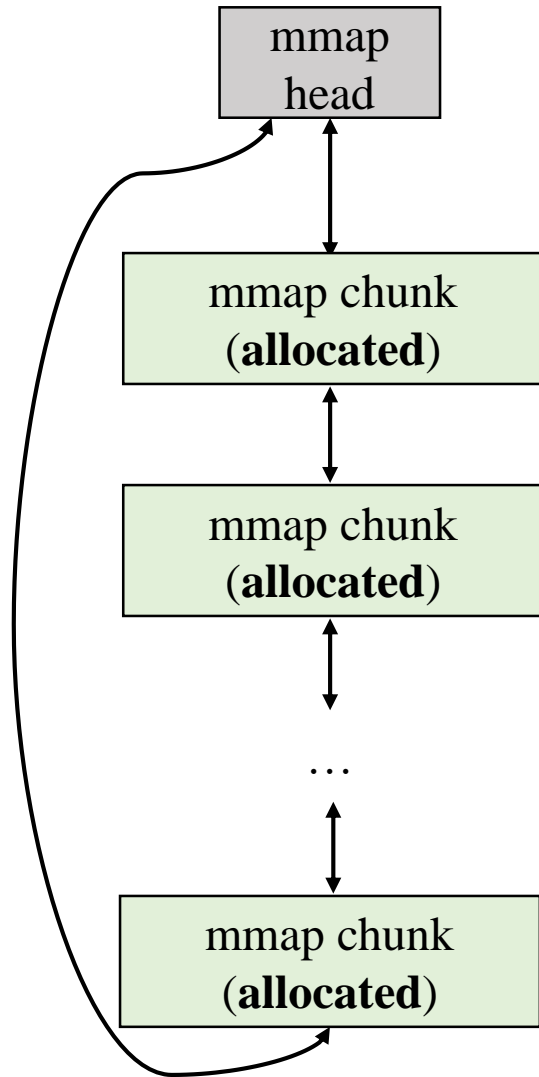
# Bins requirements

- Bins are use to manage free chunks of Heap's segment.
- A bin is a circular doubly-linked list of free chunk(s) (*next slide*)
- Add the chunk to **rear** of the bin.
- You should manage 11 bins
  - bin[0]-bin[10] hold chunks with fixed size (*next slide*)
  - Every chunk size should be **the nth power of 2** (n is a number of 5 to 15)
- Use the best fit size to select a chunk during memory allocation
  - If there are multiple chunks with the same size, **select the one with the lowest address**

# Bin example



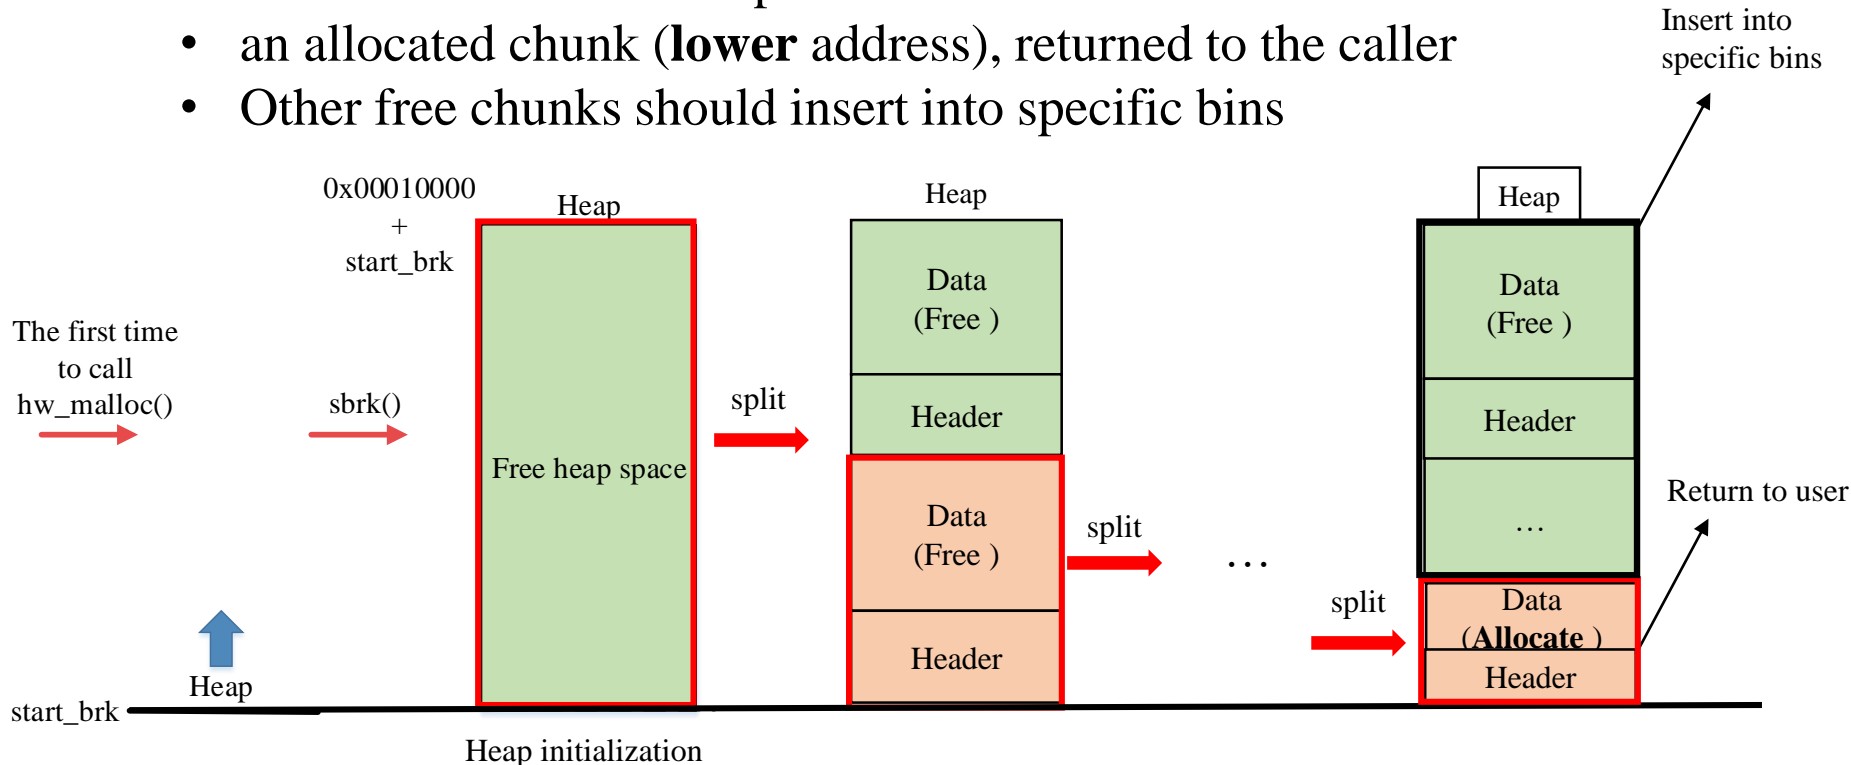
# mmap\_alloc\_list



- `mmap_alloc_list` is a circular doubly-linked list of mmap allocated chunk(s)
- mmap chunk header is same as [slide 8](#) (use `chunk_ptr_t`)
- `mmap_alloc_list` should be ordered by size (ascending)
- If there is/are multiple chunk(s) of the same size, add new chunk **after** it/them.

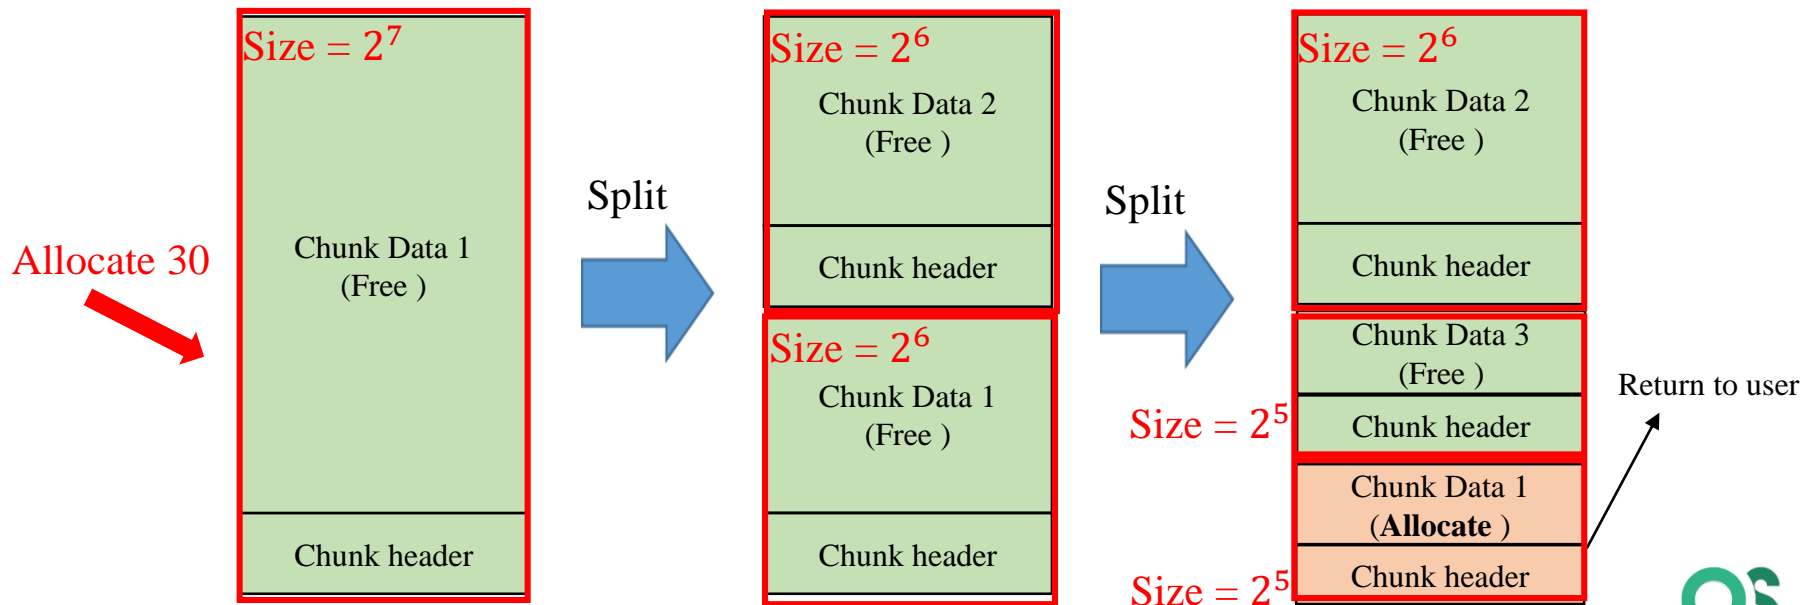
# Heap initialization & First-time Heap allocation

- Initialize the Heap before its first use :
  - Use `sbrk()` to allocate a **64KiB** heap
- After the heap initialization, Split (*next slide*) until the allocated size match the best fit of the nth power of 2
  - an allocated chunk (**lower** address), returned to the caller
  - Other free chunks should insert into specific bins



# Split

- When Heap allocation method is used, *split* may be performed :
  - If chunk size is too large (not the best fit for the allocation size), they must be split into two equal-sized chunks
  - Should split the lowest address chunk
  - Chunk size should always be the nth power of the 2
  - Must be **split until it reach the best fit for the allocation size**
- Example : (hw\_malloc(6), allocation size =30 , Chunk1 =2<sup>7</sup>)



# hw\_free( ) requirements

- Use Chunk header information (*slides 7, 8*) to check the address was allocated by mmap or Heap allocation method.

## mmap free method :

- Use **munmap()** system call to free the space
- **The chunk header should be free too.**
- This free chunk **do not need to be added into bin.**  
(bin is not used in mmap allocation method)

## Heap free method :

- Use bin (*slides 9,10*) to manage free chunks
- Should follow the rules of **Merge** (*slides 15,16*)

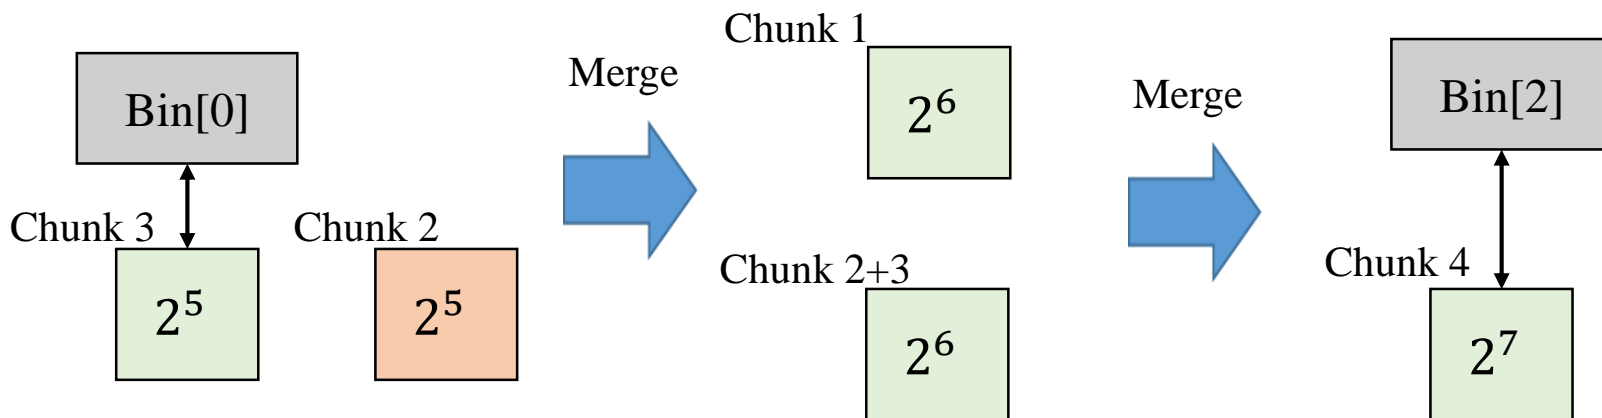
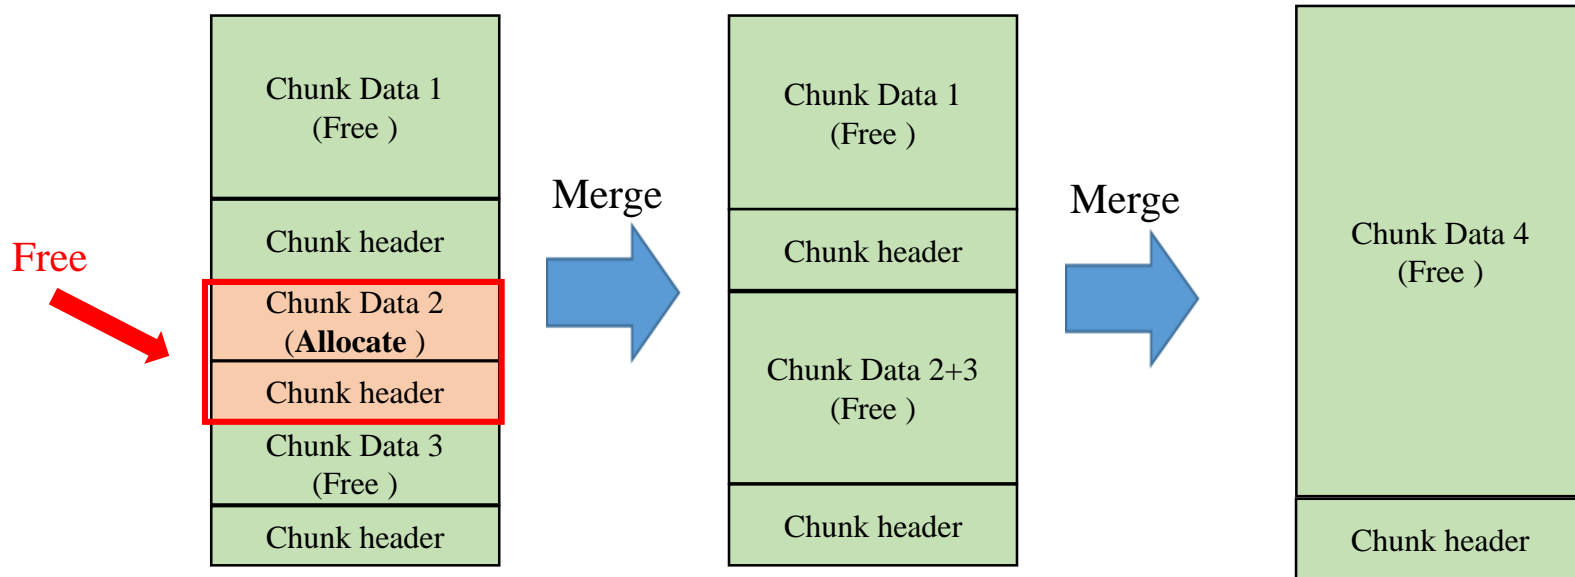
# Merge

- When Heap free method is called , adjacent free chunks must be merged into one if they have the same size :
  - **Merge operations should be repeated until the two adjacent free chunks do not have the same size**



# Merge Example

- Example: (Chunk 1 size =  $2^6$ , Chunk 2 size = Chunk 3 size =  $2^5$ )



## 2. User Application Requirements

Write a user application to test the memory allocator library

- Should receive **4 kinds of commands** :
  1. **alloc** *N*
  2. **free** *ADDR*
  3. **print** *BIN*
  4. **print** *mmap\_alloc\_list*
- Continuously receive commands from stdin until **EOF** (*Ctrl+D*)
  - Should successfully run  
“cat testfile.txt | hw4\_mm\_test > outputfile.txt”

# Commands format

## 1. **alloc N**

- Call `hw_malloc(N)` to allocate N bytes of data memory
- Print relative data address (i.e., offset between `start_brk` and the address returned by `hw_malloc()`)

## 2. **free ADDR**

- Call `hw_free()` to free the memory at (`start_brk + ADDR`)
- Print either “success” or “fail”

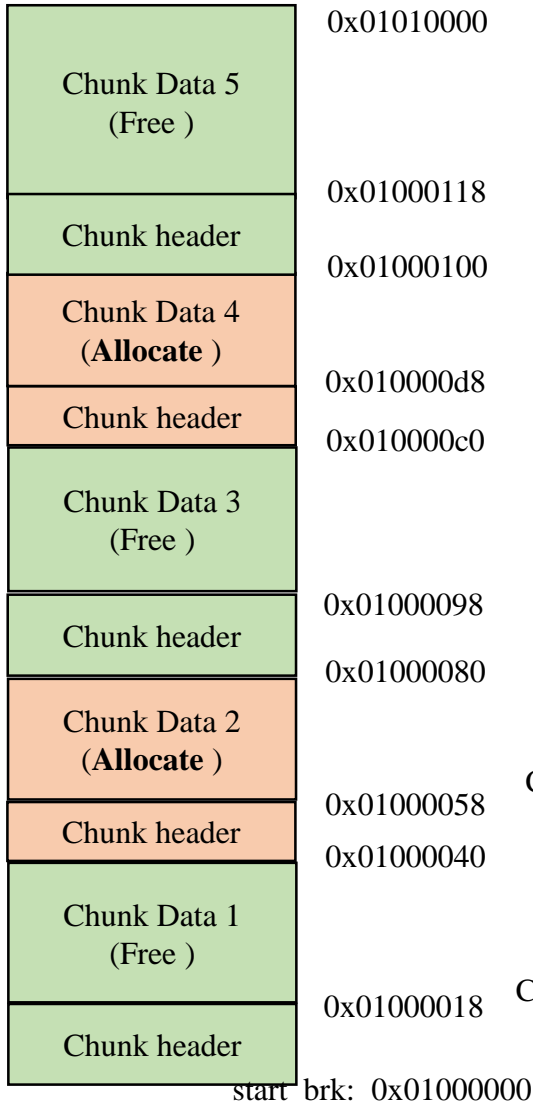
## 3. **print bin[i]**

- Print relative data address and size information of a given bin
  - `bin[i]` can be `bin[0]`, `bin[1]` ..... `bin[10]`


## 4. **print mmap\_alloc\_list**

- Print data address and size information.

# 3. Input/output format



- alloc
  - print (relative) **data** address in a line
- free
  - print success/fail in a line
- print bin[i]
  - print (relative) **chunk** address and **size** of each free chunk in the given bin, from the front to the rear
  - print a line for each chunk; pad 8 dash signs between the address and size




## Input example

```
alloc 16
alloc 16
alloc 16
alloc 16
free 0x00000018
free 0x00000098
print bin[1]
```

## Output example

```
0x00000018
0x00000058
0x000000d8
0x000000d0
success
success
0x00000000-----64
0x00000080-----64
```

# 3. Input/output format



- print `mmap_alloc_list`
  - print **chunk address** and **size** of each free chunk in the list, from the front to the rear
  - print a line for each chunk; pad 8 dash signs between the address and size


## Input example

```
alloc 33768
alloc 37864
alloc 35816
print mmap_alloc_list
```

## Output example

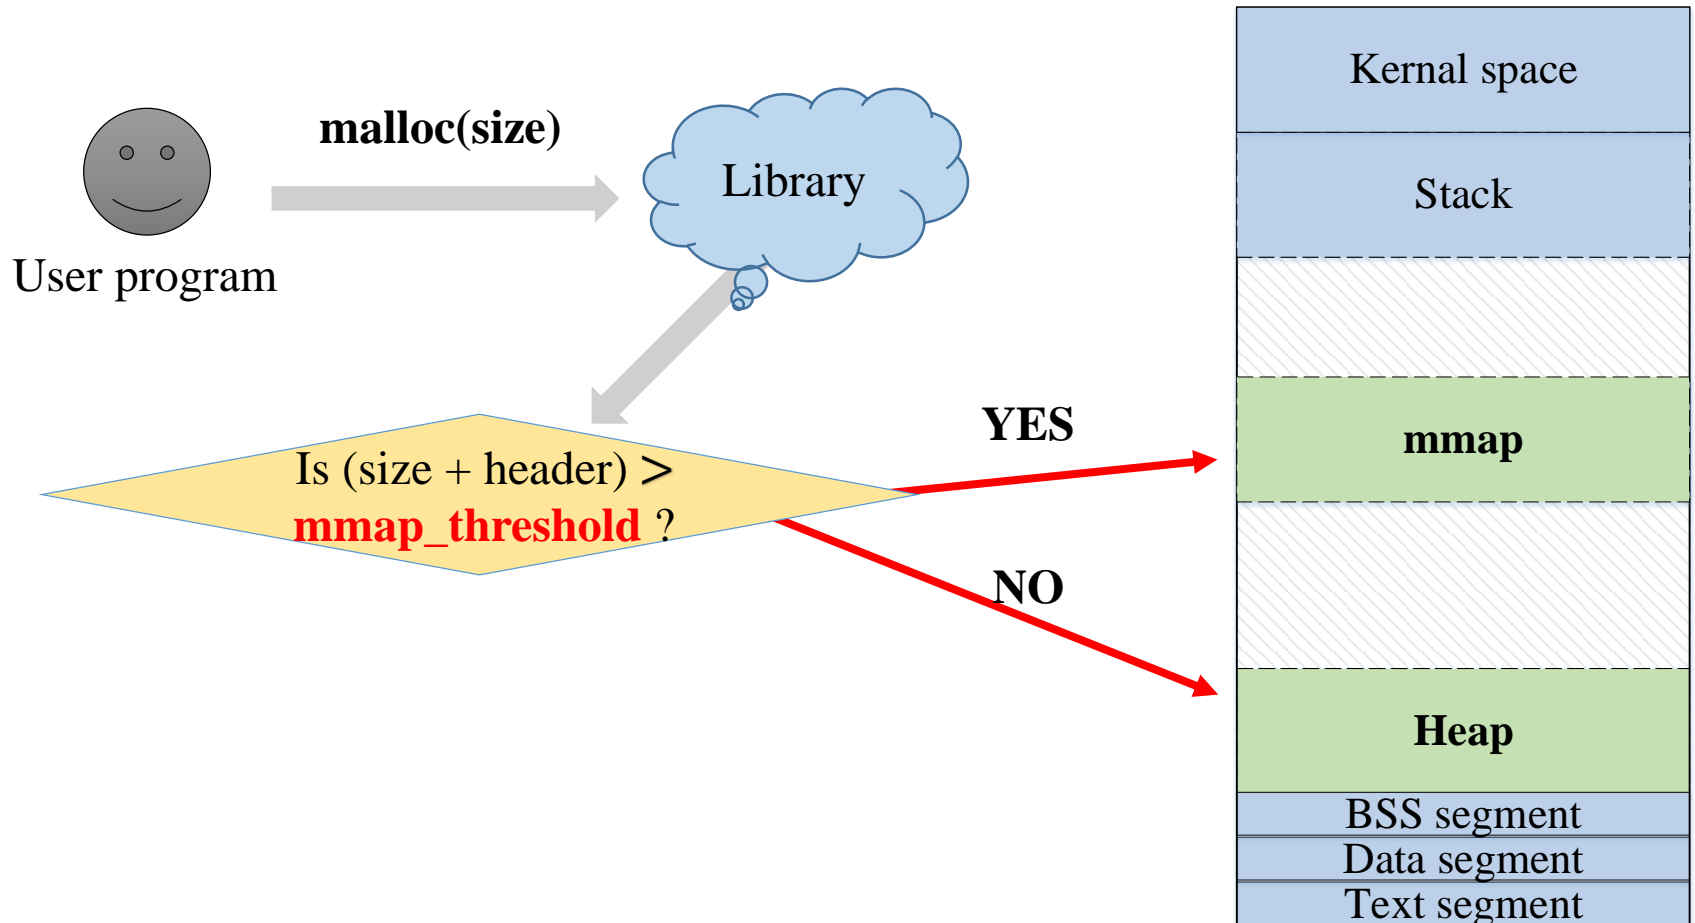
```
0xfdcf00018
0xfecff0f018
0xffcff0f018
0xfdcf00000-----33792
0xffcff0f000-----35840
0xfecff0f000-----37888
```

# Concepts (malloc / free)




# Concepts (mmap\_threshold)

- Use `mmap_threshold` to decide the memory allocate method

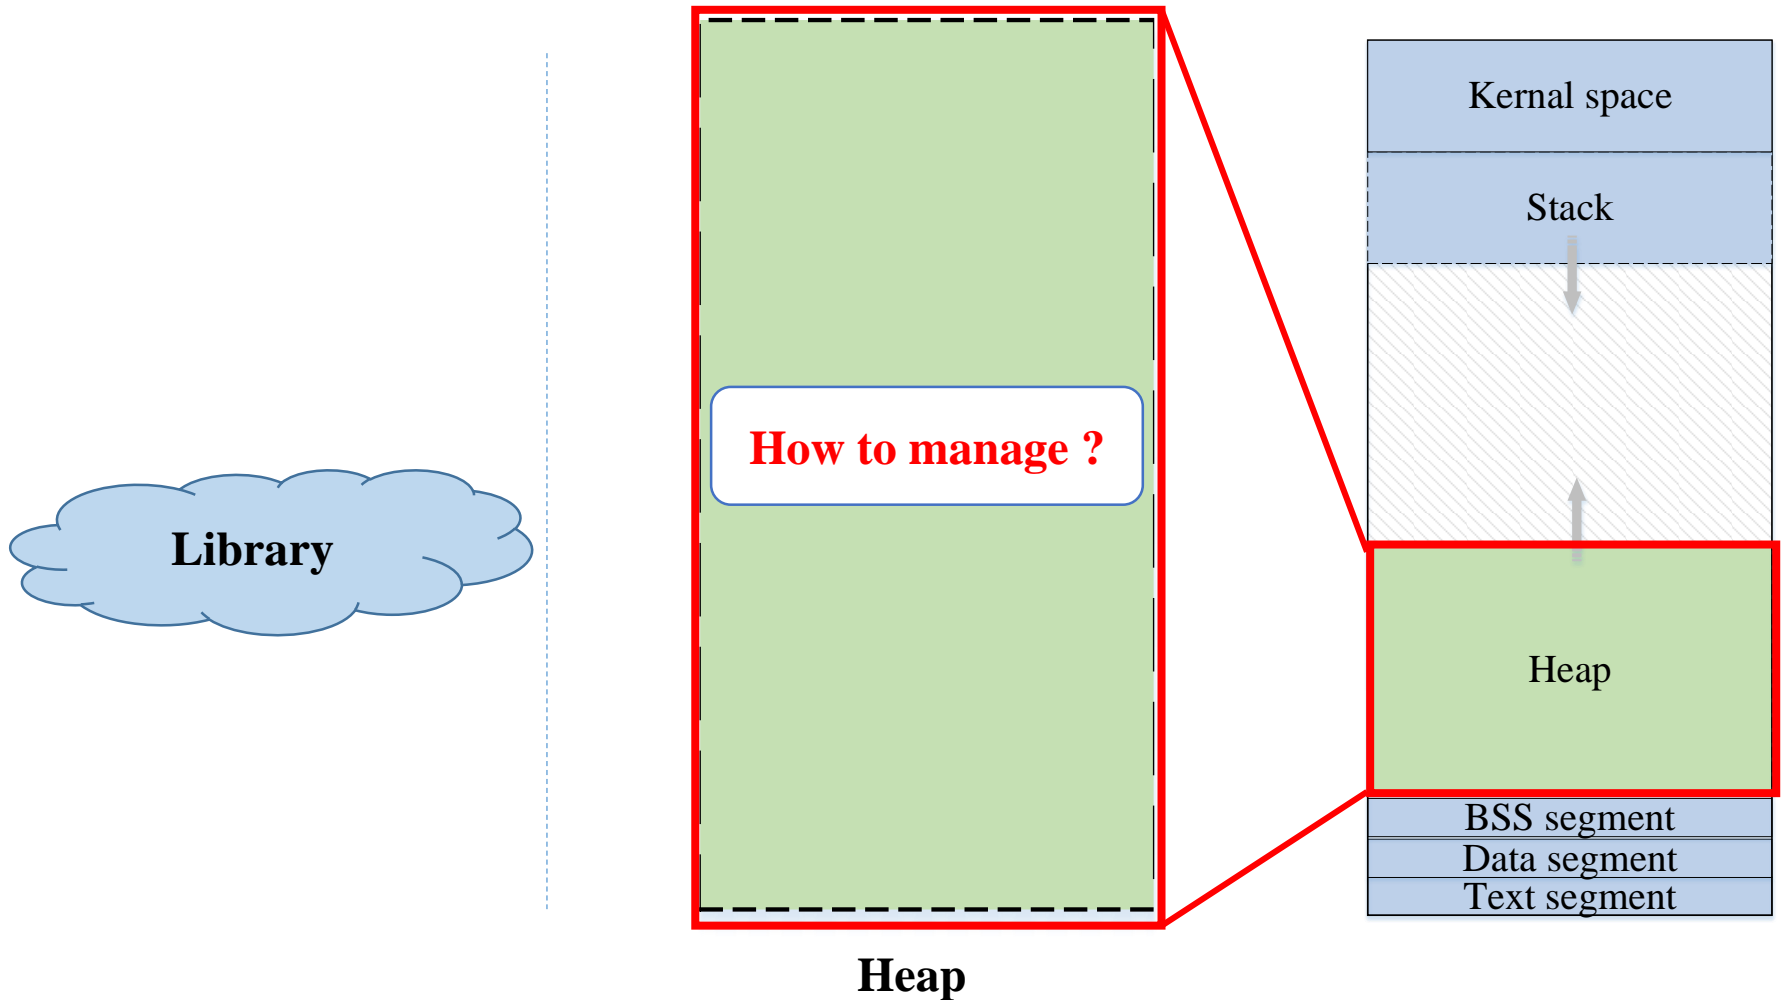


# Concepts (Heap initialization)

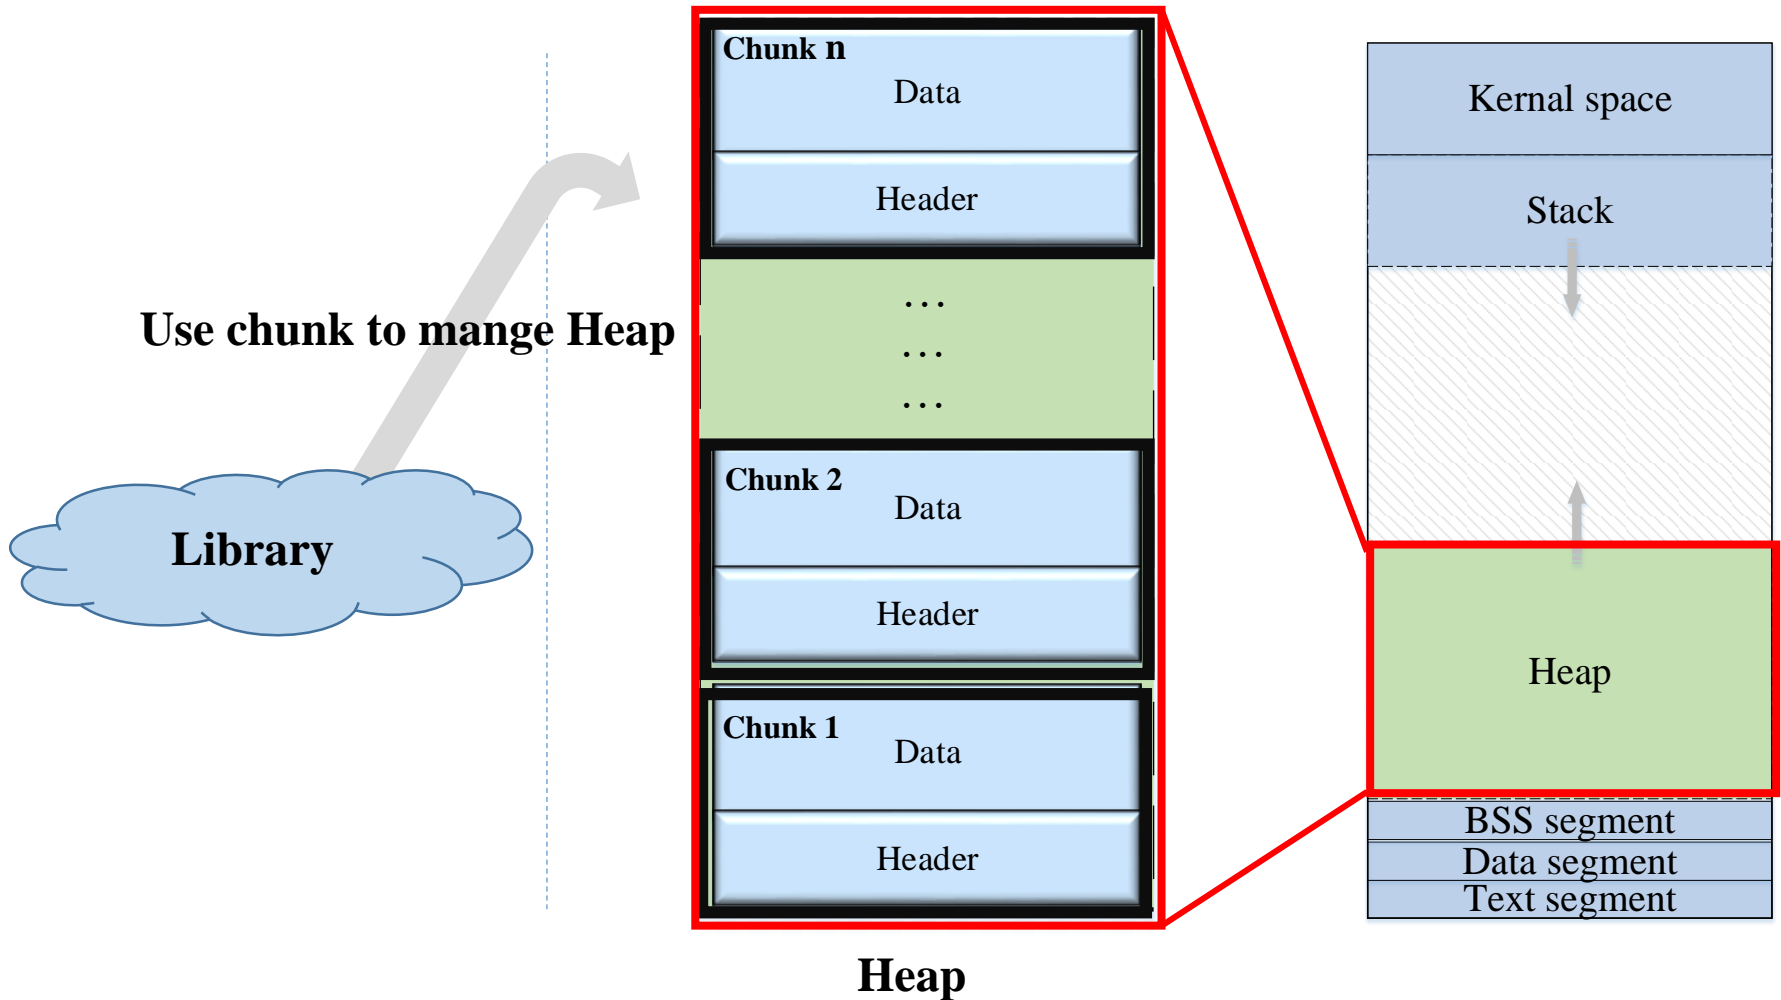





# Concepts (heap)



# Concepts (heap segment)




# Concepts (chunk)



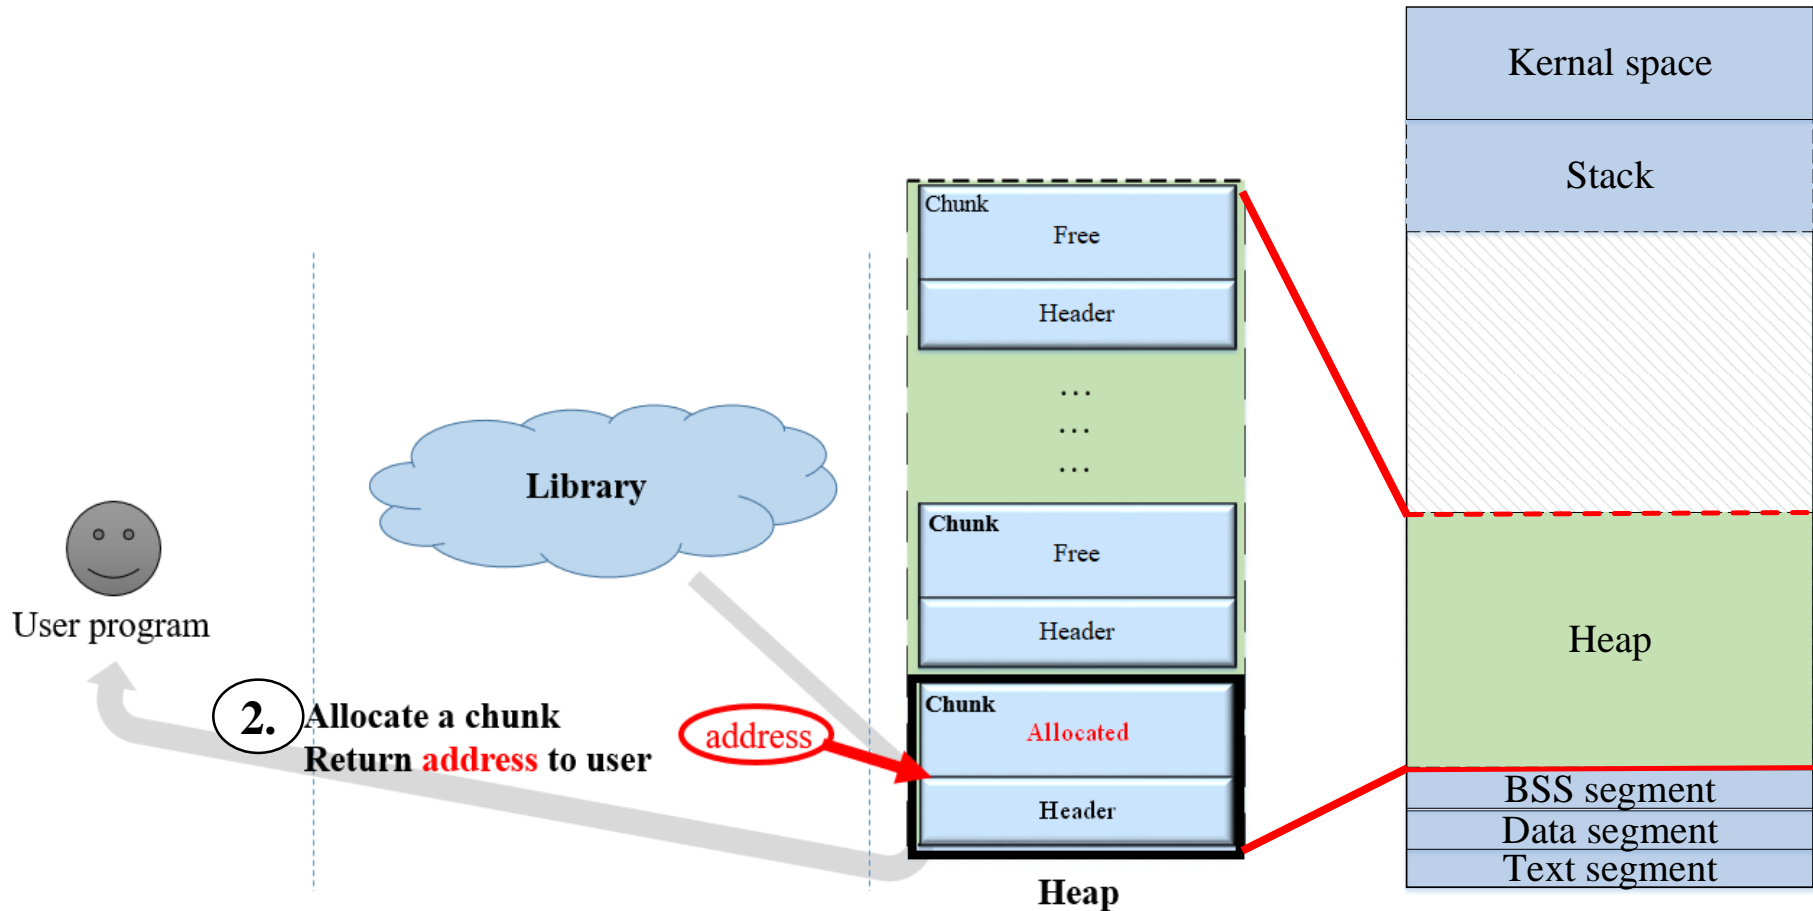
**Chunk**

- **Return address**  
(starting address of the data part)
- **Chunk information**  
(e.g: size, flag, prev, next...)

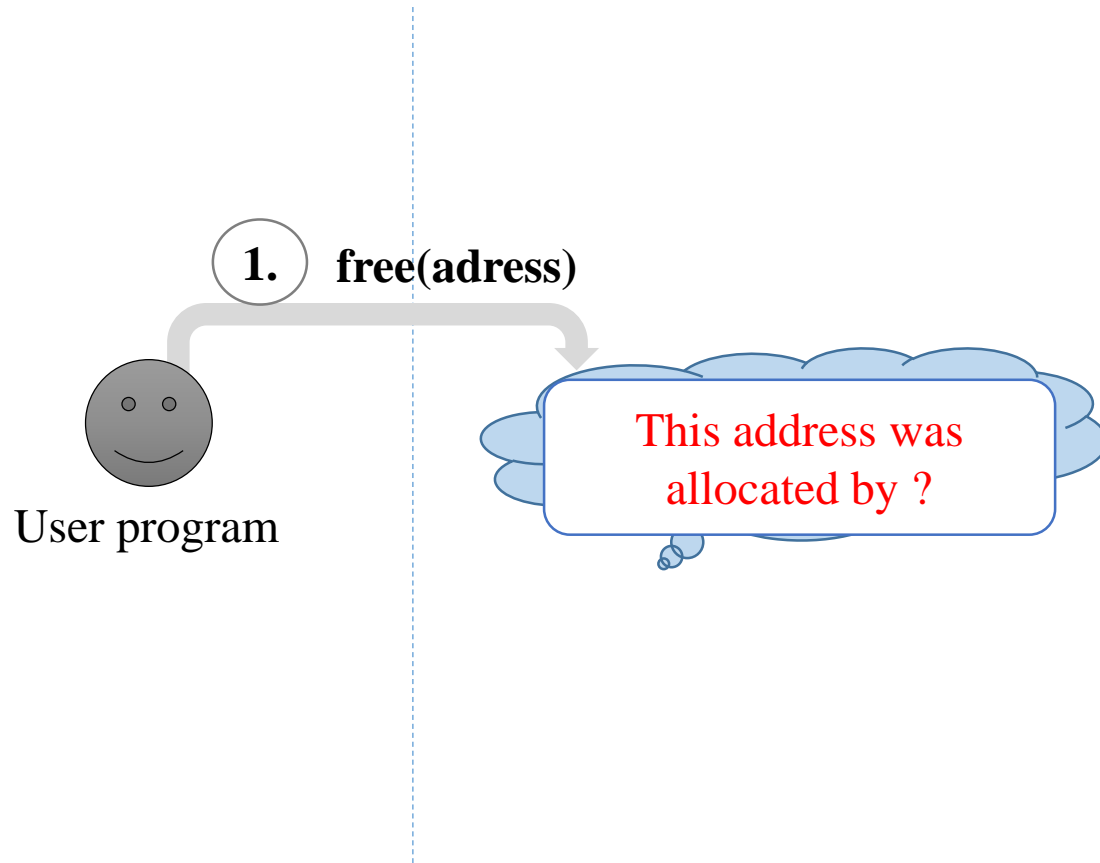
# Concepts (Heap initialization)



# Concepts (heap segment)

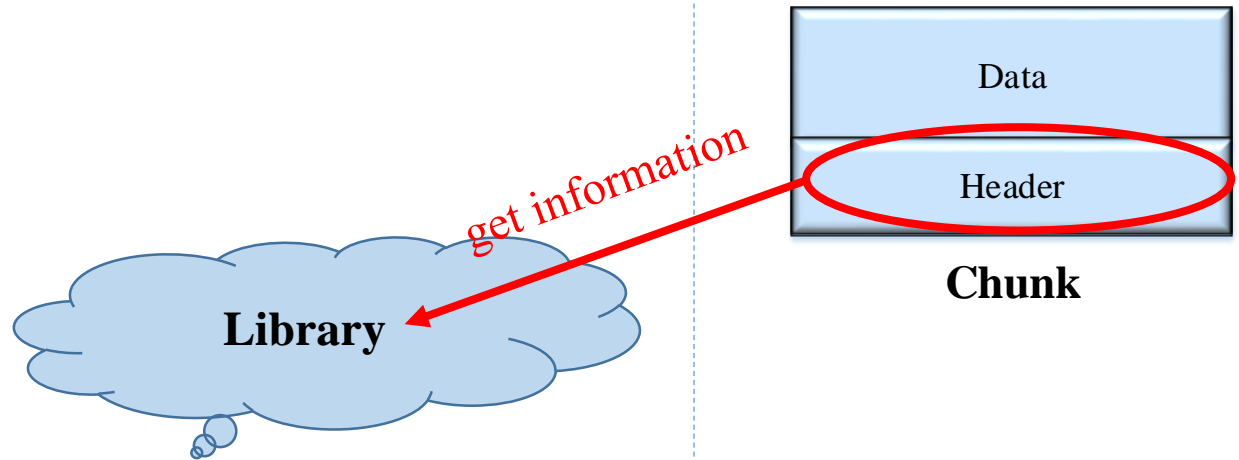



# Concepts (free)

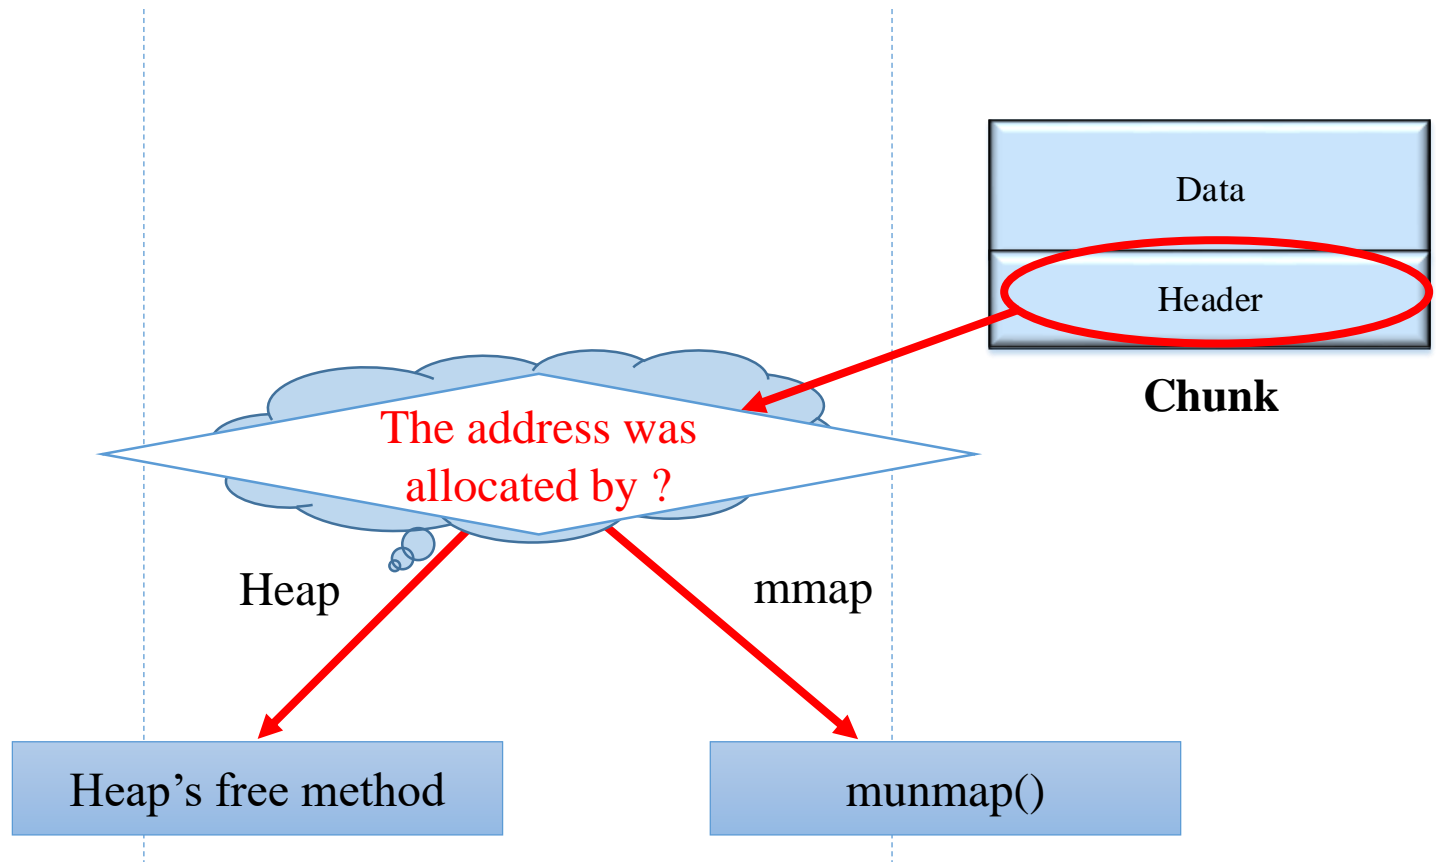


# Concepts (free)

- Get information from Chunk header to know which memory allocate method was use



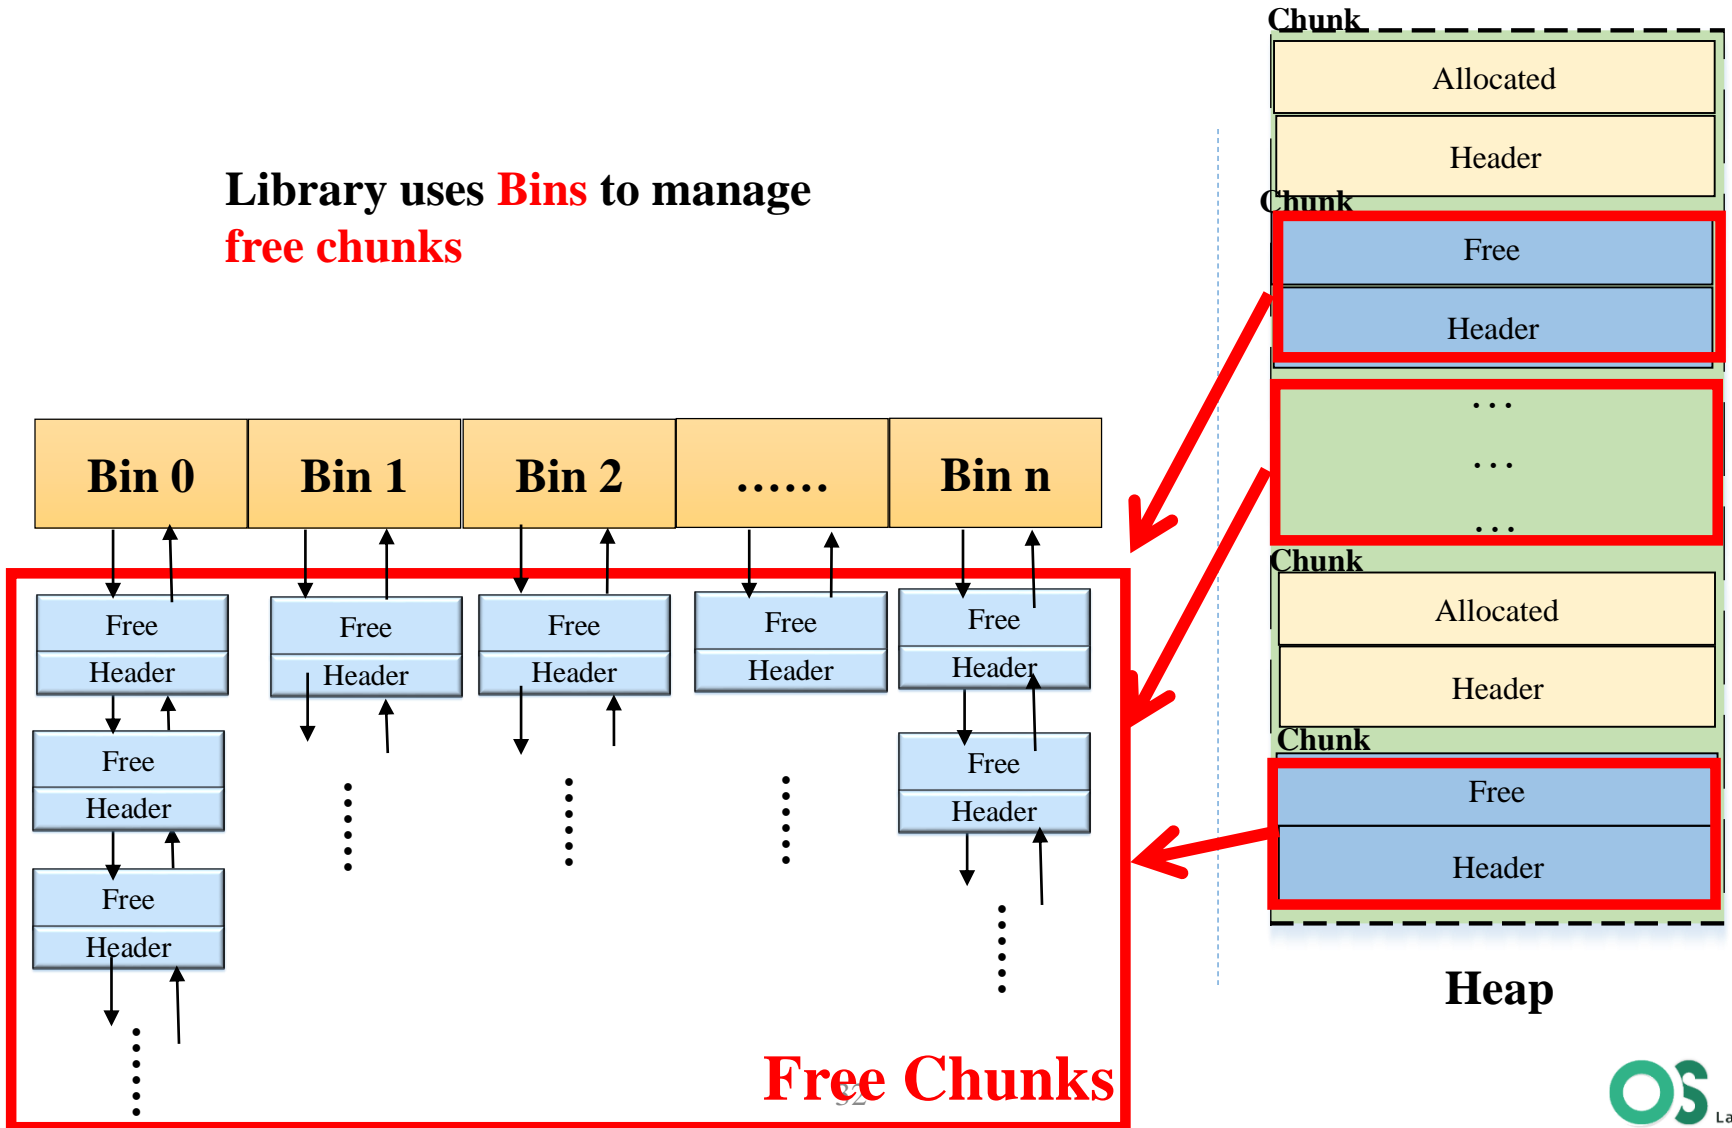
# Concepts (free)



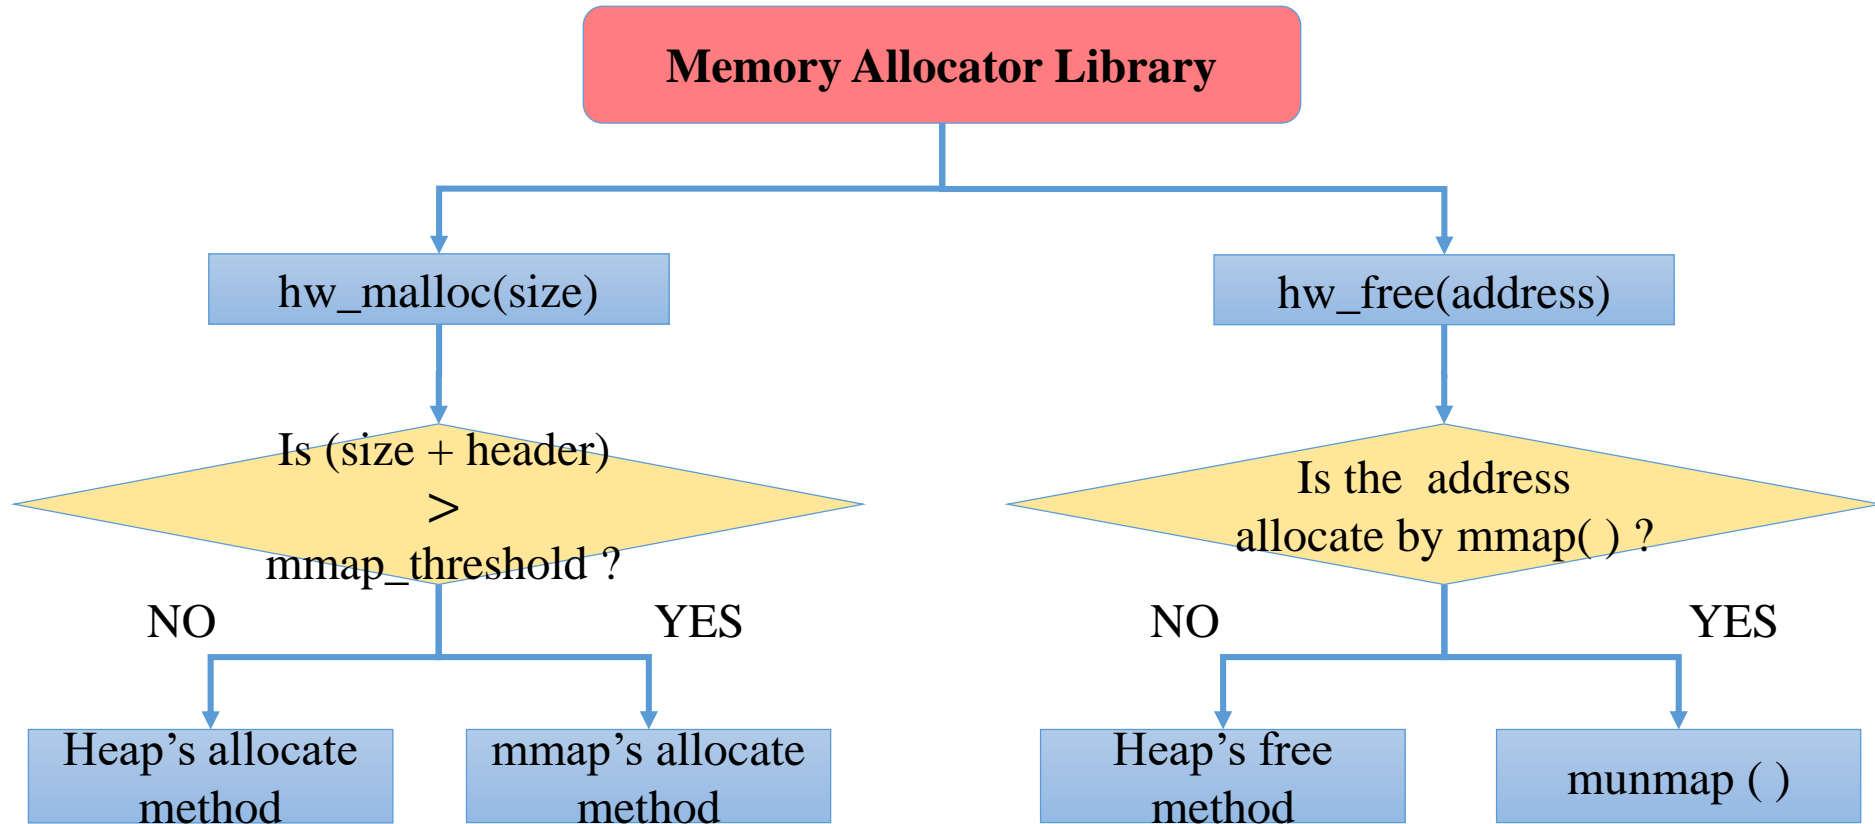


# Concepts (Heap's free method)

Library uses **Bins** to manage  
free chunks



# Concepts (Architecture)



# References

- sbrk()
  - [Linux man page](#)
- Streams, pipes, and redirects
  - [IBM](#)